

Kestrel

An XMPP-Based Framework for Many Task
Computing Applications

Lance Stout
Mike Murphy
Sebastien Goasguen

HISTORY/PURPOSE

Kestrel's Goals

Lightweight / Easy to set up

- Run cross-platform without re-compiling
- No extensive, manual configuration
- Minimal dependencies

Detect Irregular Resource Outages

- Know quickly if a worker process terminates with kill -9

Traverse NAT

High Availability / Reliability

EXTENSIBLE MESSAGING AND PRESENCE PROTOCOL (XMPP)

XMPP Benefits

Presence notifications

- Always aware of the status of the worker pool
- Always receive unavailable status updates

Indirect Communication

- All messages sent through server
- NAT and subnet traversal

Identifiers

- Address workers without knowing IP addresses
- Workers can be grouped using JIDs

Jabber IDs username@server/resource

worker@kestrel_pool/42

(Only use to group small numbers of workers)

worker42@kestrel_pool/

(One username per worker is best for large pools)

machine27@kestrel_pool/core2

Messages

Kestrel uses JSON for message contents


- Differentiated by “type” attribute
- Can be sent directly from an instant messaging client (GoogleTalk/Pidgin)

Example:


```
{“type”: “profile”,
  “os”: “Linux”,
  “ram”: 4096,
  “cores”: 4,
  “provides”: [“FOO”, “BAR”]}
```

ARCHITECTURE

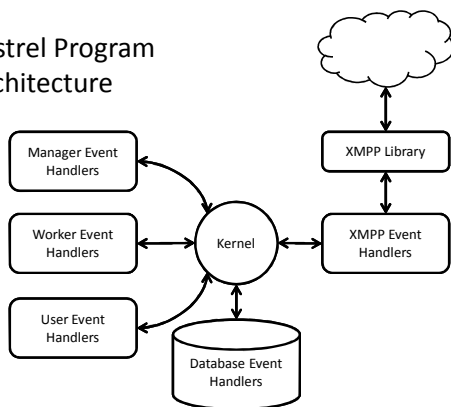
Kestrel Network Architecture (Actual)



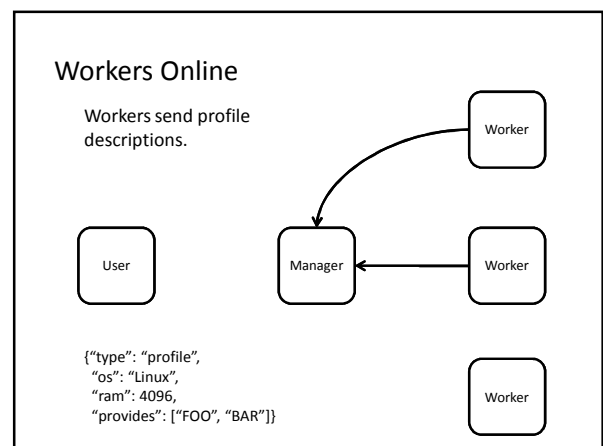
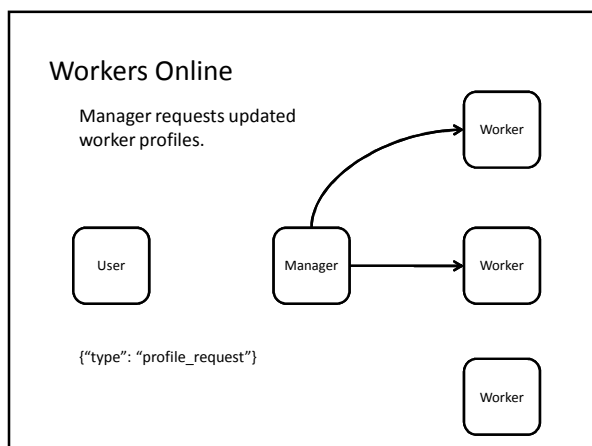
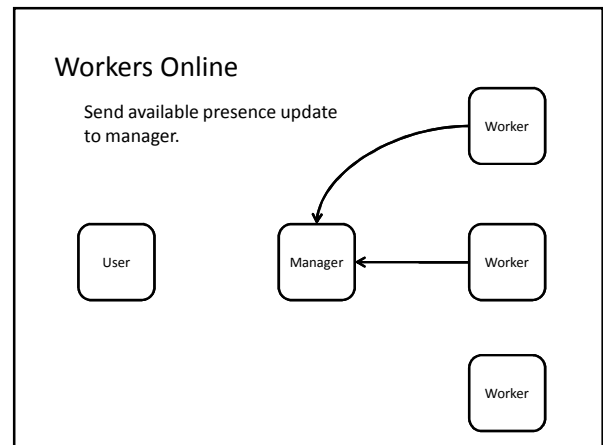
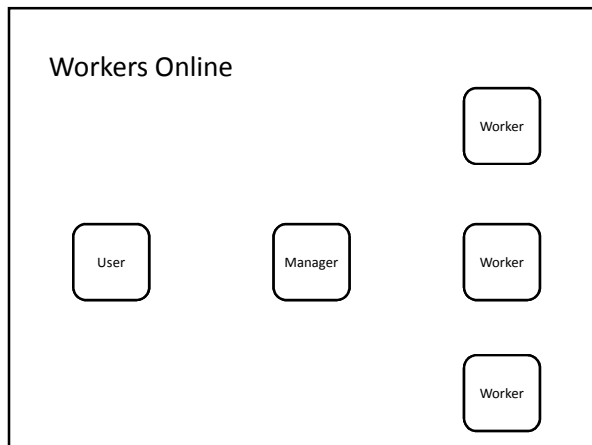
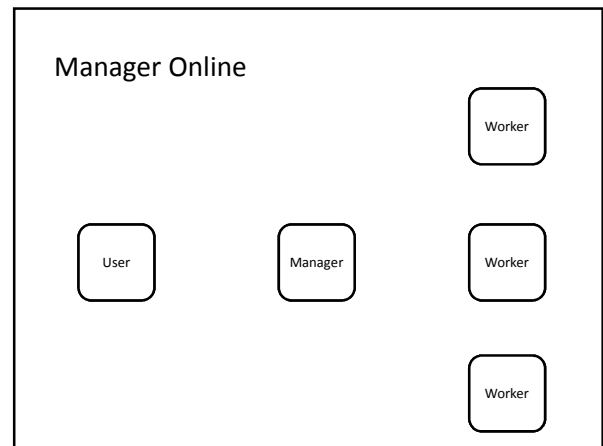
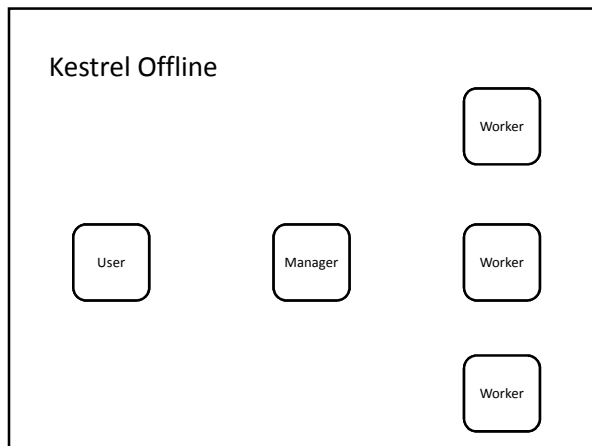
Kestrel Network Architecture (Logical)

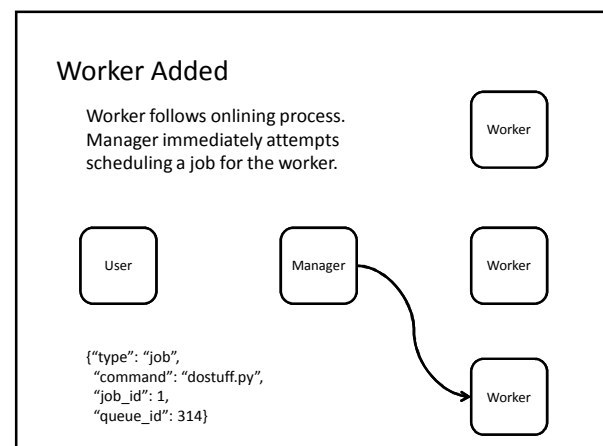
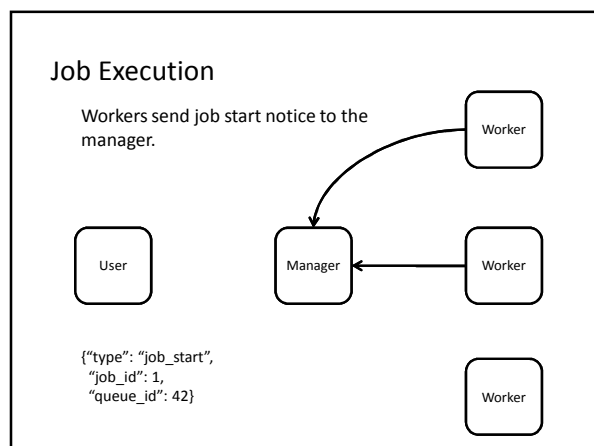
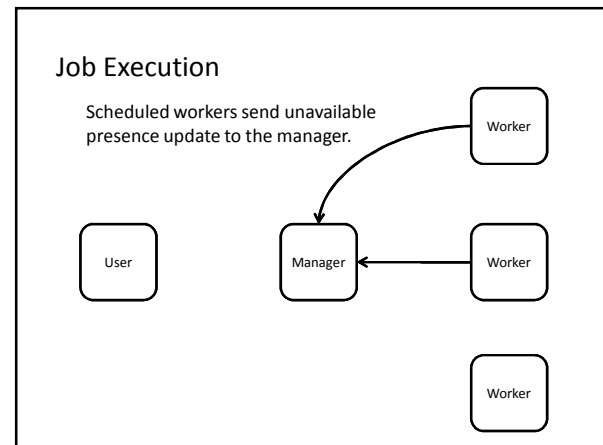
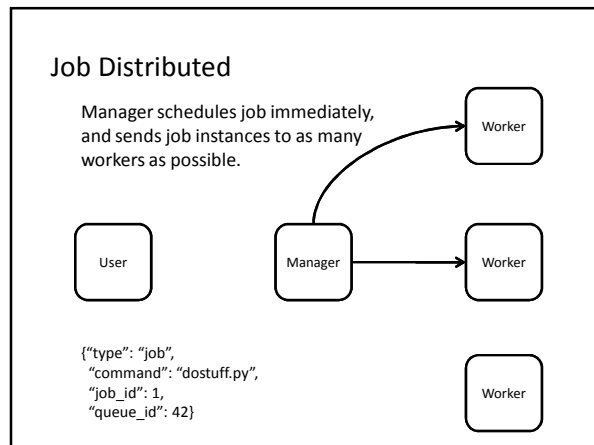
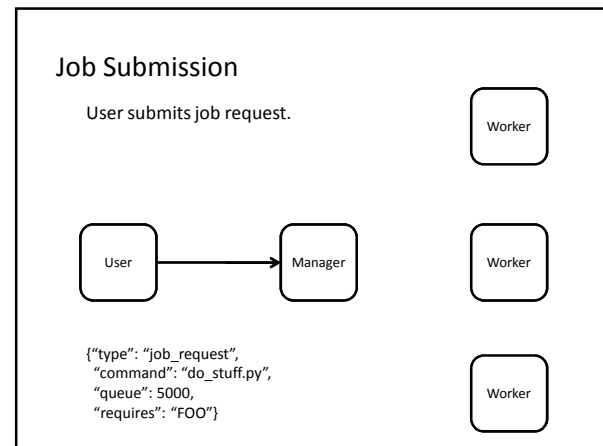
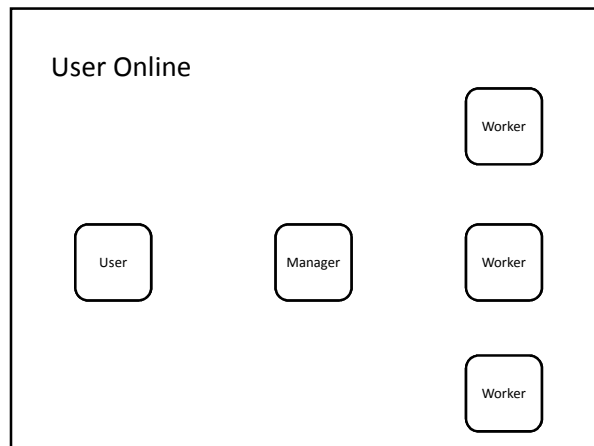


Kestrel Program Architecture




PSEUDO-DEMO






Worker Dropped

XMPP server sends unavailable presence to the manager. Any job assigned to the worker is added back to the queue.




Job Instance Finished

Workers send job finish notice to the manager.




Job Instance Finished

Workers send available presence update to the manager. Manager attempts scheduling more jobs.




Job Finished

Manager sends notice to user.



RESULTS




50,000 Jobs, 917 Workers, sleep 0

Finished in 103 seconds.


Dispatched 480 jobs per second.

Very few instances running concurrently due to short execution times.




50,000 Jobs, 917 Workers, sleep 0

Why the uptick? We're still working on that one.
Probably due to server processing faster as resources are freed.




50,000 Jobs, 910 Workers, sleep 1

Finished in 108 seconds




50,000 Jobs, 910 Workers, sleep 1

About the same as last time, including uptick.




50,000 Jobs, 914 Workers, sleep 10

Finished in 543 seconds, or 9 minutes 3 seconds.




50,000 Jobs, 914 Workers, sleep 10

Longer, uniform execution time creates stair step pattern and no uptick.



50,000 Jobs, 902 Workers, sleep random

Jobs lasted between 0 and 10 seconds.
Finished in 387 seconds, or 6 minutes 27 seconds.



50,000 Jobs, 902 Workers, sleep random

No stair step pattern or uptick this time.